

1. Which of our open ML Research Engineer roles are you interested in, in order?

I would be interested in any but the order would be Model Evaluation, General then Applied Research. Model Evaluation first because of my experience with building a 3-stage evaluation harness for an offline RL project and analyzing how the learned policies work against a fixed trivial setpoint.

2. Describe an ML idea you were initially excited about but later decided was wrong or unimportant. What changed your mind?

While working on the offline RL project, I believed that wrapping a learned policy in simple handwritten rules would be a safe way to deploy it, but mine ended up being the worst of the four policies. It was not because of the model. But it was because of a one-line mistake in my own rule. I learned that the wrapper needs to be scrutinized just like the model.

3. Describe a time when you embarked on a sidequest / took initiative outside your core job responsibilities. What value did you add to your organization? What drove you to action?

My team used to keep re-reading utility incentive program manuals by hand for every project and financial modelling of the projects. Seeing them go through manuals every time, I built an internal RAG tool over those texts (Using Fast API, React, LangChain, ChromaDB) that answers with citations to the source clause and also calculates the incentives for the project. Now the team uses it for any incentive-related questions swiftly and also uses it during client calls to check their requirements.

4. How could better accuracy fail to translate into more useful forecasts for a weather model?

A forecast model that is highly optimized on RMSE smooths its prediction toward the averages, missing the extremes. But in real life these extreme events are what is most important. These events are rare, missing one barely affects the accuracy computed over a year. But in reality these are the events the forecast models are made for.

5. Describe a time you changed how you use LLMs in your work — switching models, tools, or workflows. What did you observe that prompted the change?

I used to run one frontier model for all of my work (research, code, idea validation etc). There wasn't a single moment that changed it, I just noticed I was paying frontier prices for work a smaller model can do, and that the models are better for reviewing rather than writing the entire code. Now I split by task, and I am also building a MetaRouter that picks the coding agent from the git diff.